Deep Learning with PyTorch: A 60 Minute Blitz

Author: Soumith Chintala

Goal of this tutorial:

- Understand PyTorch's Tensor library and neural networks at a high level.
- Train a small neural network to classify images

This tutorial assumes that you have a basic familiarity of numpy

Note

Make sure you have the torch and torchvision packages installed.

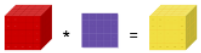

What is PyTorch?

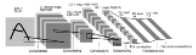

Neural Networks

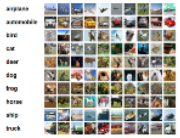

Training a classifier

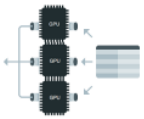

Optional: Data Parallelism

Previous

Next