

Tensors

Tensors behave almost exactly the same way in PyTorch as they do in Torch.

Create a tensor of size (5 x 7) with uninitialized memory:

```
import torch
a = torch.empty(5, 7, dtype=torch.float)
```

Initialize a double tensor randomized with a normal distribution with mean=0, var=1:

```
a = torch.randn(5, 7, dtype=torch.double)
print(a)
print(a.size())
```

Out:

```
tensor([[ 0.2095, -1.5431, -0.1285, -0.3743, -0.4776, -0.4502,  0.6270],
        [ 0.9514, -1.7134, -0.8382, -0.1200, -0.2205, -1.1100,  1.4029],
        [-0.2199,  1.7155,  1.1164, -0.2394, -0.2693, -1.9801,  0.6137],
        [-0.2756, -1.3669, -0.2064,  2.4318, -1.1294, -0.8584,  0.1801],
        [ 0.5671,  0.9556,  0.5539,  0.0953,  0.6315, -1.2639,  0.3836]])
torch.Size([5, 7])
```

Note

`torch.Size` is in fact a tuple, so it supports the same operations

Inplace / Out-of-place

The first difference is that ALL operations on the tensor that operate in-place on it will have an `_inplace` postfix. For example, `add_` is the out-of-place version, and `add_.` is the in-place version.

```
a = fill_(3.5)
# a has now been filled with the value 3.5

b = a.add(4.0)
# a is still filled with 3.5
# new tensor b is returned with values 3.5 + 4.0 = 7.5

print(a, b)
```

Out:

```
tensor([[ 3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000],
        [ 3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000],
        [ 3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000],
        [ 3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000],
        [ 3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000,  3.5000]])
tensor([[ 7.5000,  7.5000,  7.5000,  7.5000,  7.5000,  7.5000,  7.5000],
        [ 7.5000,  7.5000,  7.5000,  7.5000,  7.5000,  7.5000,  7.5000],
        [ 7.5000,  7.5000,  7.5000,  7.5000,  7.5000,  7.5000,  7.5000],
        [ 7.5000,  7.5000,  7.5000,  7.5000,  7.5000,  7.5000,  7.5000],
        [ 7.5000,  7.5000,  7.5000,  7.5000,  7.5000,  7.5000,  7.5000]])
```

Some operations like `narrow` do not have in-place versions, and hence, `narrow_` does not exist. Similarly, some operations like `fill_` do not have an out-of-place version, so `fill` does not exist.

Zero Indexing

Another difference is that Tensors are zero-indexed. (In lua, tensors are one-indexed)

```
b = a[0, 3] # select 1st row, 4th column from a
```

Tensors can be also indexed with Python's slicing

```
b = a[:, 3:5] # selects all rows, 4th column and 5th column from a
```

No camel casing

The next small difference is that all functions are now NOT camelCase anymore. For example

`indexAdd` is now called `index_add_`.

```
x = torch.ones(5, 5)
print(x)
```

Out:

```
tensor([[ 1.,  1.,  1.,  1.,  1.],
        [ 1.,  1.,  1.,  1.,  1.],
        [ 1.,  1.,  1.,  1.,  1.],
        [ 1.,  1.,  1.,  1.,  1.],
        [ 1.,  1.,  1.,  1.,  1.]])
```

```
z = torch.empty(5, 2)
z[:, 0] = 10
z[:, 1] = 100
print(z)
```

Out:

```
tensor([[ 10., 100.],
        [ 10., 100.],
        [ 10., 100.],
        [ 10., 100.],
        [ 10., 100.]])
```

```
x.index_add_(1, torch.tensor([4, 0], dtype=torch.long), z)
print(x)
```

Out:

```
tensor([[ 101.,  1.,  1.,  1.,  11.],
        [ 101.,  1.,  1.,  1.,  11.],
        [ 101.,  1.,  1.,  1.,  11.],
        [ 101.,  1.,  1.,  1.,  11.],
        [ 101.,  1.,  1.,  1.,  11.]])
```

Numpy Bridge

Converting a torch Tensor to a numpy array and vice versa is a breeze. The torch Tensor and numpy array will share their underlying memory locations, and changing one will change the other.

Converting torch Tensor to numpy Array

```
a = torch.ones(5)
print(a)
```

Out:

```
tensor([ 1.,  1.,  1.,  1.,  1.])
```

```
b = a.numpy()
print(b)
```

Out:

```
[ 1.  1.  1.  1.  1.]
```

```
a.add_(1)
print(a)
print(b) # see how the numpy array changed in value
```

Out:

```
tensor([ 2.,  2.,  2.,  2.,  2.])
[ 2.  2.  2.  2.  2.]
```

Converting numpy Array to torch Tensor

```
import numpy as np
a = np.ones(5)
b = torch.from_numpy(a)
np.add(a, 1, out=a)
print(a)
print(b) # see how changing the np array changed the torch Tensor automatically
```

Out:

```
[ 2.  2.  2.  2.  2.]
tensor([ 2.,  2.,  2.,  2.,  2.], dtype=torch.float64)
```

All the Tensors on the CPU except a `CharTensor` support converting to NumPy and back.

CUDA Tensors

CUDA Tensors are nice and easy in pytorch, and transferring a CUDA tensor from the CPU to GPU will retain its underlying type.

```
# let us run this cell only if CUDA is available
if torch.cuda.is_available():
    # creates a LongTensor and transfers it
    # to GPU as torch.cuda.LongTensor
    a = torch.full((10, 1, 3), device=torch.device("cuda"))
    print(a)
    b = a.to(torch.device("cpu"))
    # transfers it to CPU, back to
    # being a torch.LongTensor
```

Out:

```
<class 'torch.Tensor'>
```

Total running time of the script: (0 minutes 0.004 seconds)

Download Python source code: tensor_tutorial.py

Download Jupyter notebook: tensor_tutorial.ipynb

Gallery generated by Sphinx Gallery

Previous

Next