

## BEGINNER TUTORIALS

[Deep Learning with PyTorch: A 60 Minute Blitz](#)[PyTorch for former Torch users](#)[Learning PyTorch with Examples](#)[Tensors](#)[Autograd](#)[nn module](#)[Examples](#)[Tensors](#)[Autograd](#)[PyTorch: Tensors and autograd](#)[PyTorch: Defining new autograd functions](#)[TensorFlow: Static Graphs](#)[nn module](#)[Transfer Learning tutorial](#)[Data Loading and Processing Tutorial](#)[Deep Learning for NLP with Pytorch](#)

## INTERMEDIATE TUTORIALS

[Classifying Names with a Character-Level RNN](#)[Generating Names with a Character-Level RNN](#)[Translation with a Sequence to Sequence Network and Attention](#)[Reinforcement Learning \(DQN\) tutorial](#)

## PyTorch: Tensors and autograd

A fully-connected ReLU network with one hidden layer and no biases, trained to predict  $y$  from  $x$  by minimizing squared Euclidean distance.

This implementation computes the forward pass using operations on PyTorch Tensors, and uses PyTorch autograd to compute gradients.

A PyTorch Tensor represents a node in a computational graph. If `x` is a Tensor that has `x.requires_grad=True` then `x.grad` is another Tensor holding the gradient of `x` with respect to some scalar value.

```
import torch

dtype = torch.float
device = torch.device("cpu")
# dtype = torch.device("cuda:0") # Uncomment this to run on GPU

# N is batch size; D_in is input dimension;
# H is hidden dimension; D_out is output dimension.
N, D_in, H, D_out = 64, 1000, 100, 10

# Create random Tensors to hold input and outputs.
# Setting requires_grad=False indicates that we do not need to compute gradients
# with respect to these Tensors during the backward pass.
x = torch.randn(N, D_in, device=device, dtype=dtype)
y = torch.randn(N, D_out, device=device, dtype=dtype)

# Create random Tensors for weights.
# Setting requires_grad=True indicates that we want to compute gradients with
# respect to these Tensors during the backward pass.
w1 = torch.randn(D_in, H, device=device, dtype=dtype, requires_grad=True)
w2 = torch.randn(H, D_out, device=device, dtype=dtype, requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    # Forward pass: compute predicted y using operations on Tensors; these
    # are exactly the same operations we used to compute the forward pass using
    # Tensors, but we do not need to keep references to intermediate values since
    # we are not implementing the backward pass by hand.
    y_pred = x.mm(w1).clamp(min=0).mm(w2)

    # Compute and print loss using operations on Tensors.
    # Now loss is a Tensor of shape (1,)
    # loss.item() gets the a scalar value held in the loss.
    loss = (y_pred - y).pow(2).sum()
    print(t, loss.item())

    # Use autograd to compute the backward pass. This call will compute the
    # gradient of loss with respect to all Tensors with requires_grad=True.
    # After this call w1.grad and w2.grad will be Tensors holding the gradient
    # of the loss with respect to w1 and w2 respectively.
    loss.backward()

    # Manually update weights using gradient descent. Wrap in torch.no_grad()
    # because weights have requires_grad=True, but we don't need to track this
    # in autograd.
    # An alternative way is to operate on weight.data and weight.grad.data.
    # Recall that tensor.data gives a tensor that shares the storage with
    # the tensor, but doesn't track history.
    # You can also use torch.optim.SGD to achieve this.
    with torch.no_grad():
        w1 -= learning_rate * w1.grad
        w2 -= learning_rate * w2.grad

    # Manually zero the gradients after updating weights
    w1.grad.zero_()
    w2.grad.zero_()
```

Total running time of the script: ( 0 minutes 0.000 seconds)

[Download Python source  
code: two\\_layer\\_net\\_autograd.py](#)[Download Jupyter notebook:  
two\\_layer\\_net\\_autograd.ipynb](#)

Gallery generated by Sphinx-Gallery

[Previous](#)[Next](#)