

Sequence Models and Long-Short Term Memory Networks

At this point, we have seen various feed-forward networks. That is, there is no state maintained by the network at all. This might not be the behavior we want. Sequence models are central to NLP: they are models where there is some sort of dependence through time between your inputs. The classical example of a sequence model is the Hidden Markov Model for part-of-speech tagging. Another example is the conditional random field.

A recurrent neural network is a network that maintains some kind of state. For example, its output could be used as part of the next input, so that information can propagate along as the network passes over the sequence. In the case of an LSTM, for each element in the sequence, there is a corresponding hidden state h_t , which in principle can contain information from arbitrary points earlier in the sequence. We can use the hidden state to predict words in a language model, part-of-speech tag, and a myriad of other things.

LSTM's in Pytorch

Before getting to the example, note a few things. Pytorch's LSTM expects all of its inputs to be 3D tensors. The semantics of the axes of these tensors is important. The first axis is the sequence itself, the second indexes instances in the mini batch, and the third indexes elements of the input. We haven't discussed mini-batching, so lets just ignore that and assume we will always have just 1 dimension on the second axis. If we want to run the sequence model over the sentence "The cow jumped", our input should look like

$$\begin{bmatrix} \text{one vector} \\ \vec{q}_{\text{The}} \\ \vec{q}_{\text{cow}} \\ \vec{q}_{\text{jumped}} \end{bmatrix}$$

Except remember there is an additional 2nd dimension with size 1.

In addition, you could go through the sequence one at a time, in which case the 1st axis will have size 1 also.

Let's see a quick example.

```
# Author: Robert Dufresne

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

torch.manual_seed(1)

# ...

lstm = nn.LSTM(3, 3) # Input dim is 3, output dim is 3
inputs = [torch.randn(1, 3) for _ in range(5)] # make a sequence of length 5
# initialize the hidden state
hidden = (torch.randn(1, 3, 3), torch.randn(1, 3, 3))
for i in inputs:
    # Step through the sequence one element at a time.
    # after each step, hidden contains the hidden state.
    out, hidden = lstm(i.view(1, 1, -1), hidden)

# alternatively, we can do the entire sequence all at once.
# the first value returned by LSTM is all of the hidden states throughout
# the sequence. the second is just the most recent hidden state
# compare the last state of "out" with "hidden" below; they are the same!
# the reason for this is that:
# "out" will give you access to all hidden states in the sequence
# "hidden" will allow you to continue the sequence and backpropagate,
# by passing it as an argument to the lstm at a later time
# add the extra 2nd dimension
inputs = torch.cat(inputs).view(len(inputs), 1, -1)
hidden = (torch.randn(1, 3, 3), torch.randn(1, 3, 3)) # clean out hidden state
out, hidden = lstm(inputs, hidden)
print(out)

# ...

Out:
tensor([[[[-0.0187,  0.3713, -0.2944]],
         [[-0.3521,  0.3826, -0.2973]],
         [[-0.3191,  0.6781, -0.1957]],
         [[-0.1626,  0.6941, -0.1627]],
         [[-0.3368,  0.6959, -0.6526]]]])
        (tensor([[[[-0.3368,  0.6959, -0.6526]]]), tensor([[[[-0.9825,  0.4715, -0.9633]]]]))
```

Example: An LSTM for Part-of-Speech Tagging

In this section, we will use an LSTM to get part of speech tags. We will not use Viterbi or Forward-Backward or anything like that, but as a (challenging) exercise to the reader, think about how Viterbi could be used after you have seen what is going on.

The model is as follows: let our input sentence be w_1, \dots, w_N , where $w_i \in V$, our vocab. Also, let T be our tag set, and y_i the tag of word w_i . Denote our prediction of the tag of word w_i by \hat{y}_i .

This is a structure prediction model, where our output is a sequence $\hat{y}_1, \dots, \hat{y}_N$, where $\hat{y}_i \in T$.

To do the prediction, pass an LSTM over the sentence. Denote the hidden state at timestep i as h_i . Also, assign each tag a unique index (like how we had word to_{ix} in the word embeddings section). Then our prediction rule for \hat{y}_i is

$$\hat{y}_i = \operatorname{argmax}_j (\log \operatorname{Softmax}(Ah_i + b_j))$$

That is, take the log softmax of the affine map of the hidden state, and the predicted tag is the tag that has the maximum value in this vector. Note this implies immediately that the dimensionality of the target space of A is $|T|$.

Prepare data:

```
def prepare_sequence(seq, to_ix):
    len_s = len(seq)
    for x in seq:
        return torch.tensor(ids, dtype=torch.long)

training_data = [
    ("The dog ate the apple",split()), ("DET", "NN", "n", "DET", "NN"),
    ("Everybody read that book",split()), ("N", "V", "DET", "NN")
]

word to ix = {}
for sent, tags in training_data:
    if word not in word to ix:
        word to ix[word] = len(word to ix)
print(word to ix)
tag to ix = {"DET": 0, "NN": 1, "V": 2}

# These will usually be more like 32 or 64 dimensional.
# we will keep them small, so we can see how the weights change as we train.
EMBEDDING_DIM = 6
HIDDEN_DIM = 6

# ...

Out:
[["The": 0, "dog": 1, "ate": 2, "that": 3, "apple": 4, "Everybody": 5, "read": 6, "that": 7, "I": 8]]
```

Create the model:

```
class LSTMTagger(nn.Module):

    def __init__(self, embedding_dim, hidden_dim, vocab_size, tagset_size):
        super(LSTMTagger, self).__init__()
        self.embedding_dim = embedding_dim
        self.word_embeddings = nn.Embedding(vocab_size, embedding_dim)

        # the LSTM takes word embeddings as inputs, and outputs hidden states
        # with dimensionality hidden_dim
        self.lstm = nn.LSTM(embedding_dim, hidden_dim)

        # the linear layer that maps from hidden state space to tag space
        self.linear = nn.Linear(hidden_dim, tagset_size)
        self.hidden = self._init_hidden()

    def __init_hidden(self):
        # Before we've done anything, we don't have any hidden state.
        # Refer to the Pytorch documentation to see exactly
        # why they have this dimensionality.
        # the two tensors are (num layers, minibatch size, hidden dim)
        return (torch.zeros(1, self.hidden_dim),
                torch.zeros(1, self.hidden_dim))

    def forward(self, sentence):
        embeds = self.word_embeddings(sentence)
        lstm_out, self.hidden = self.lstm(
            embeds.view(len(sentence), 1, -1), self.hidden)
        tag_space = self.linear(lstm_out.view(len(sentence), -1))
        tag_scores = F.log_softmax(tag_space, dim=1)
        return tag_scores
```

Train the model:

```
model = LSTMTagger(EMBEDDING_DIM, HIDDEN_DIM, len(word to ix), len(tag to ix))
loss function = nn.NLLLoss()
optimizer = optim.SGD(model.parameters(), lr=0.1)

# See what the scores are before training
# note that element i,j of the output is the score for tag j for word i.
# there we don't need to train, so the code is wrapped in torch.no_grad()
with torch.no_grad():
    inputs = prepare_sequence(training_data[0][0], word to ix)
    tag_scores = model(inputs)
    print(tag_scores)

for epoch in range(100): # again, normally you would NOT do 100 epochs, it is too fast
    for sentence, tags in training_data:
        # Step 1. Remember that Pytorch accumulates gradients.
        # we need to clear them out before each instance
        model.zero_grad()

        # Also, we need to clear out the hidden state of the LSTM,
        # decoupling it from its history on the last instance
        model.hidden = model._init_hidden()

        # Step 2. Get our inputs ready for the network, that is, turn them into
        # Tensors of word indices.
        sentence_in = prepare_sequence(sentence, word to ix)
        targets = prepare_sequence(tags, tag to ix)

        # Step 3. Run our forward pass.
        tag_scores = model(sentence_in)

        # Step 4. Compute the loss, gradients, and update the parameters by
        # calling optimizer.step()
        loss = loss_function(tag_scores, targets)
        loss.backward()
        optimizer.step()

# See what the scores are after training
with torch.no_grad():
    inputs = prepare_sequence(training_data[0][0], word to ix)
    tag_scores = model(inputs)

    # The sentence is "the dog ate the apple". 4 corresponds to score for tag 4
    # for word 1. The predicted tag is the maximum scoring tag.
    # Here, we can see the predicted sequence below is 0 2 2 0 2
    # since 0 is index of the maximum value of row 1,
    # 1 is the index of maximum value of the 2, etc
    # which is NOT what we got for ANN, the correct sequence!
    print(tag_scores)
```

Out:

```
tensor([[-1.1389, -1.2624, -0.6893],
        [-1.1865, -1.2396, -0.9034],
        [-1.1286, -1.2893, -0.9726],
        [-1.1396, -1.1666, -0.9619],
        [-1.8137, -1.2842, -0.8981]],
        tensor([[-0.8824, -2.8555, -0.9578],
                [-5.2313, -0.8234, -0.8314],
                [-5.8666, -0.1275, -0.8568],
                [-0.8337, -4.7869, -1.5662],
                [-5.8176, -0.8185, -4.1875]]])
```

Exercise: Augmenting the LSTM part-of-speech tagger with character-level features

In the example above, each word had an embedding, which served as the inputs to our sequence model. Let's augment the word embeddings with a representation derived from the characters of the word. We expect that this should help significantly, since character-level information like affixes have a large bearing on part-of-speech. For example, words with the affix -ly are almost always tagged as adverbs in English.

To do this, let c_w be the character-level representation of word w . Let x_w be the word embedding as before. Then the input to our sequence model is the concatenation of x_w and c_w . So if x_w has dimension 5, and c_w dimension 3, then our LSTM should accept an input of dimension 8.

To get the character level representation, do an LSTM over the characters of a word, and let c_w be the final hidden state of this LSTM. Hints:

- There are going to be two LSTM's in your new model. The original one that outputs POS tag scores, and the new one that outputs a character-level representation of each word.
- To do a sequence model over characters, you will have to embed characters. The character embeddings will be the input to the character LSTM.

Total running time of the script: (0 minutes 0.984 seconds)

Download Python source code: sequence_model1_tutorial.py

Download Jupyter notebook: sequence_model1_tutorial.ipynb

Gallery generated by Sphinx Gallery

Previous

Next

©Copyright 2017, PyTorch.

Built with [Sphinx](#) using a theme provided by [Read the Docs](#).